

# hw3\_ChengjunGuo

Chengjun Guo

February 2023

## 1 SGD+

Equation:

$$\begin{aligned}v_{t+1} &= \mu * v_t + g_{t+1} \\ p_{t+1} &= p_t - \text{lr} * v_{t+1}\end{aligned}$$

In the equations, p means the parameter and t means the previous iteration.  $\mu$  is the momentum coefficient that inherit the step of previous step. g is the gradient of loss. Compared to SGD, this method is basically inheriting the speed in previous step. With this algorithm, the learning rate would be more adaptive than original SGD algorithm and it will converge faster.

## 2 Adam

Equations:

$$\begin{aligned}m_{t+1} &= \beta_1 * m_t + (1 - \beta_1) * g_{t+1} \\ v_{t+1} &= \beta_2 * v_t + (1 - \beta_2) * (g_{t+1})^2 \\ p_{t+1} &= p_t - \text{lr} * \frac{\hat{m}}{\sqrt{\hat{v}_{t+1} + \epsilon}}\end{aligned}$$

where


$$\begin{aligned}\hat{m}_k &= \frac{m_k}{1 - \beta_1^k} \\ \hat{v}_k &= \frac{v_k}{1 - \beta_2^k}\end{aligned}$$

$\beta_1$  and  $\beta_2$  are the user defined variables. Adam is combining the momentum based logic and the sparse gradient into a single algorithm. Here's the motion of development from adagrad to rmsprop. The adagrad is developed based on the intuition that whenever a partial derivative becomes non-zero, the rareness of such occurrences could mean that those dimensions carry high class discriminatory information and it should take larger steps. Then adagrad runs into problem that the monotonically increasing value for the denominator could case the learning rate for a parameter to become vanishing small. RMSprop replace the summation in denominator with its average over training iterations to fix it.

Then adam comes to use the momentum and be adaptive to different component of gradient.

### 3 Plots

#### 3.1 one neuron




(a) 1e-3



(b) 5e-5

Figure 1: one neuron

#### 3.2 multi neuron



(a) 1e-3



(b) 5e-5

Figure 2: multi neuron

## 4 Discussion

Based on the training loss plots, we can find that one neuron is more stabilize than multi neuron. When the learning rate is  $1e-3$ , it can be seen that Adam converge faster than sgd+ and sgd+ converge faster than sgd. Adam is increasing after it achieves a far lower loss than the sgd+. When the learning rate is  $5e-5$ , sgd is not converging and sgd+ converge faster at faster and adam converge faster after that. Adam and sgd+ is better at handling low learning rate with the momentum.

## 5 Code

One neuron:

---

```
#!/usr/bin/env python

## one_neuron_classifier.py

"""
A one-neuron model is characterized by a single
expression that you see in the value
supplied for the constructor parameter "expressions". In
the expression supplied, the
names that begin with 'x' are the input variables and the
names that begin with the
other letters of the alphabet are the learnable
parameters.
"""

import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
import sys
sys.path.append("E:\ECE60146DL\hw3_new\
ComputationalGraphPrimer-1.1.2\
ComputationalGraphPrimer")

import sys, os, os.path
import numpy as np
import re
import operator
import math
import random
import torch
from collections import deque
```

```

import copy
import matplotlib.pyplot as plt
import networkx as nx

seed = 0
random.seed(seed)
np.random.seed(seed)

from ComputationalGraphPrimer import *

class ComputationalGraphPrimerPlus(
    ComputationalGraphPrimer):
    def run_training_loop_one_neuron_model(self,
        training_data, momentum_coe):
        """
        The training loop must first initialize the
        learnable parameters. Remember, these are the
        symbolic names in your input expressions for the
        neural layer that do not begin with the
        letter 'x'. In this case, we are initializing
        with random numbers from a uniform
        distribution
        over the interval (0,1).
        """
        self.vals_for_learnable_params = {param: random.
            uniform(0, 1) for param in self.
            learnable_params}
        self.gamma = momentum_coe
        self.bias = random.uniform(0, 1) ## Adding the
            bias improves class discrimination.
        self.prev_grad = {param: 0 for param in self.
            learnable_params}
        self.prev_bias = 0
        ## We initialize it to a random number.

    class DataLoader:
        """
        To understand the logic of the dataloader, it
        would help if you first understand how
        the training dataset is created. Search for
        the following function in this file:

            gen_training_data(self)

```

As you will see in the implementation code for this method, the training dataset consists of a Python dict with two keys, 0 and 1, the former points to a list of all Class 0 samples and the latter to a list of all Class 1 samples. In each list, the data samples are drawn from a multi-dimensional Gaussian distribution. The two classes have different means and variances. The dimensionality of each data sample is set by the number of nodes in the input layer of the neural network.

The data loader's job is to construct a batch of samples drawn randomly from the two lists mentioned above. And it must also associate the class label with each sample separately.

```

def __init__(self, training_data, batch_size)
:
    self.training_data = training_data
    self.batch_size = batch_size
    self.class_0_samples = [(item, 0) for
        item in
            self.
                training_data
                [0]] ##
                Associate
                label 0 with
                each sample
    self.class_1_samples = [(item, 1) for
        item in
            self.
                training_data
                [1]] ##
                Associate
                label 1 with
                each sample

def __len__(self):
    return len(self.training_data[0]) + len(
        self.training_data[1])

```

```

def _getitem(self):
    cointoss = random.choice([0, 1])  ## When
    a batch is created by getbatch(), we
    want the
    ## samples to be chosen randomly from
    the two lists
    if cointoss == 0:
        return random.choice(self.
                               class_0_samples)
    else:
        return random.choice(self.
                               class_1_samples)

def getbatch(self):
    batch_data, batch_labels = [], []  ##
    First list for samples, the second for
    labels
    maxval = 0.0  ## For approximate batch
    data normalization
    for _ in range(self.batch_size):
        item = self._getitem()
        if np.max(item[0]) > maxval:
            maxval = np.max(item[0])
        batch_data.append(item[0])
        batch_labels.append(item[1])
    batch_data = [item / maxval for item in
                  batch_data]  ## Normalize batch data
    batch = [batch_data, batch_labels]
    return batch

data_loader = DataLoader(training_data,
                          batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_iterations = 0.0  ## Average the
loss over iterations for printing out
## every N iterations during the training loop
.
for i in range(self.training_iterations):
    data = data_loader.getbatch()
    data_tuples = data[0]
    class_labels = data[1]
    y_preds, deriv_sigmoids = self.
        forward_prop_one_neuron_model(data_tuples)
    ## FORWARD PROP of data

```

```

        loss = sum([(abs(class_labels[i] - y_preds[i]
                        )) ** 2 for i in range(len(class_labels))
                    ]) ## Find loss
        loss_avg = loss / float(len(class_labels))
        ## Average the loss over batch
        avg_loss_over_iterations += loss_avg
        if i % (self.display_loss_how_often) == 0:
            avg_loss_over_iterations /= self.
                display_loss_how_often
            loss_running_record.append(
                avg_loss_over_iterations)
            print("[iter=%d]  loss = %.4f" % (i + 1,
                avg_loss_over_iterations)) ## Display
                average loss
            avg_loss_over_iterations = 0.0 ## Re-
                initialize avg loss
        y_errors = list(map(operator.sub,
            class_labels, y_preds))
        y_error_avg = sum(y_errors) / float(len(
            class_labels))
        deriv_sigmoid_avg = sum(deriv_sigmoids) /
            float(len(class_labels))
        data_tuple_avg = [sum(x) for x in zip(*
            data_tuples)]
        data_tuple_avg = list(map(operator.truediv,
            data_tuple_avg,
                                [float(len(
                                    class_labels))]
                                * len(
                                    class_labels)))

        self.
            backprop_and_update_params_one_neuron_model
            (y_error_avg, data_tuple_avg,
            deriv_sigmoid_avg) ## BACKPROP loss
    return loss_running_record
# plt.figure()
# plt.plot(loss_running_record)
# plt.show()

def backprop_and_update_params_one_neuron_model(self,
    y_error, vals_for_input_vars, deriv_sigmoid):
    """
    As should be evident from the syntax used in the
    following call to backprop function,

```

```

self.
    backprop_and_update_params_one_neuron_model
    ( y_error_avg, data_tuple_avg,
      deriv_sigmoid_avg)

```

^^^

^^^

^^^

the values fed to the backprop function for its three arguments are averaged over the training samples in the batch. This in keeping with the spirit of SGD that calls for averaging the information retained in the forward propagation over the samples in a batch.

See Slide 59 of my Week 3 slides for the math of back propagation for the One-Neuron network.

```

"""
input_vars = self.independent_vars
input_vars_to_param_map = self.var_to_var_param[
    self.output_vars[0]]
param_to_vars_map = {param: var for var, param in
    input_vars_to_param_map.items()}
vals_for_input_vars_dict = dict(zip(input_vars,
    list(vals_for_input_vars)))
vals_for_learnable_params = self.
    vals_for_learnable_params
for i, param in enumerate(self.
    vals_for_learnable_params):
    ## Calculate the next step in the parameter
    hyperplane
    #          step = self.learning_rate *
    y_error * vals_for_input_vars_dict[
    input_vars[i]] * deriv_sigmoid
grad = y_error * vals_for_input_vars_dict[
    param_to_vars_map[param]] * deriv_sigmoid
step = self.learning_rate * grad + self.gamma
    * self.prev_grad[param]
    ## Update the learnable parameters
    self.prev_grad[param] = step
    self.vals_for_learnable_params[param] += step
grad = y_error * deriv_sigmoid

```



```

self.prev_bias = self.gamma * self.prev_bias +
    self.learning_rate * grad
self.bias += self.prev_bias

```

```

class ComputationalGraphPrimerAdam(
    ComputationalGraphPrimer):
    def run_training_loop_one_neuron_model(self,
        training_data, beta1, beta2):
        """
        The training loop must first initialize the
        learnable parameters. Remember, these are the
        symbolic names in your input expressions for the
        neural layer that do not begin with the
        letter 'x'. In this case, we are initializing
        with random numbers from a uniform
        distribution
        over the interval (0,1).
        """
        self.vals_for_learnable_params = {param: random.
            uniform(0, 1) for param in self.
            learnable_params}
        self.epsilon = 1e-8
        self.beta1 = beta1
        self.beta2 = beta2
        self.bias = random.uniform(0, 1) ## Adding the
            bias improves class discrimination.
        self.prev_m = {param: 0 for param in self.
            learnable_params}
        self.prev_v = {param: 0 for param in self.
            learnable_params}
        self.prev_biasm = 0
        self.prev_biasv = 0
        ## We initialize it to a random number.

class DataLoader:
    """
    To understand the logic of the dataloader, it
    would help if you first understand how
    the training dataset is created. Search for
    the following function in this file:

        gen_training_data(self)
    """

```

As you will see in the implementation code for this method, the training dataset consists of a Python dict with two keys, 0 and 1, the former points to a list of all Class 0 samples and the latter to a list of all Class 1 samples. In each list, the data samples are drawn from a multi-dimensional Gaussian distribution. The two classes have different means and variances. The dimensionality of each data sample is set by the number of nodes in the input layer of the neural network.

The data loader's job is to construct a batch of samples drawn randomly from the two lists mentioned above. And it must also associate the class label with each sample separately.

```
def __init__(self, training_data, batch_size):
    :
    self.training_data = training_data
    self.batch_size = batch_size
    self.class_0_samples = [(item, 0) for
        item in
            self.
                training_data
                [0]] ##
                Associate
                label 0 with
                each sample
    self.class_1_samples = [(item, 1) for
        item in
            self.
                training_data
                [1]] ##
                Associate
                label 1 with
                each sample

def __len__(self):
    return len(self.training_data[0]) + len(
        self.training_data[1])
```

```

def _getitem(self):
    cointoss = random.choice([0, 1])  ## When
    a batch is created by getbatch(), we
    want the
    ## samples to be chosen randomly from
    the two lists
    if cointoss == 0:
        return random.choice(self.
                               class_0_samples)
    else:
        return random.choice(self.
                               class_1_samples)

def getbatch(self):
    batch_data, batch_labels = [], []  ##
    First list for samples, the second for
    labels
    maxval = 0.0  ## For approximate batch
    data normalization
    for _ in range(self.batch_size):
        item = self._getitem()
        if np.max(item[0]) > maxval:
            maxval = np.max(item[0])
        batch_data.append(item[0])
        batch_labels.append(item[1])
    batch_data = [item / maxval for item in
                  batch_data]  ## Normalize batch data
    batch = [batch_data, batch_labels]
    return batch

data_loader = DataLoader(training_data,
                          batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_iterations = 0.0  ## Average the
loss over iterations for printing out
## every N iterations during the training loop
.
for i in range(self.training_iterations):
    data = data_loader.getbatch()
    data_tuples = data[0]
    class_labels = data[1]
    y_preds, deriv_sigmoids = self.
        forward_prop_one_neuron_model(data_tuples)
    ## FORWARD PROP of data

```

```

        loss = sum([(abs(class_labels[i] - y_preds[i]
                        )) ** 2 for i in range(len(class_labels))
                    ]) ## Find loss
        loss_avg = loss / float(len(class_labels))
        ## Average the loss over batch
        avg_loss_over_iterations += loss_avg
        if i % (self.display_loss_how_often) == 0:
            avg_loss_over_iterations /= self.
                display_loss_how_often
            loss_running_record.append(
                avg_loss_over_iterations)
            print("[iter=%d]  loss = %.4f" % (i + 1,
                avg_loss_over_iterations)) ## Display
                average loss
            avg_loss_over_iterations = 0.0 ## Re-
                initialize avg loss
        y_errors = list(map(operator.sub,
            class_labels, y_preds))
        y_error_avg = sum(y_errors) / float(len(
            class_labels))
        deriv_sigmoid_avg = sum(deriv_sigmoids) /
            float(len(class_labels))
        data_tuple_avg = [sum(x) for x in zip(*
            data_tuples)]
        data_tuple_avg = list(map(operator.truediv,
            data_tuple_avg,
                                [float(len(
                                    class_labels))]
                                * len(
                                    class_labels)))

        self.
            backprop_and_update_params_one_neuron_model
            (y_error_avg, data_tuple_avg,
            deriv_sigmoid_avg) ## BACKPROP loss
    return loss_running_record
# plt.figure()
# plt.plot(loss_running_record)
# plt.show()

def backprop_and_update_params_one_neuron_model(self,
    y_error, vals_for_input_vars, deriv_sigmoid):
    """

```

As should be evident from the syntax used in the following call to backprop function,

```

self.
    backprop_and_update_params_one_neuron_model
    ( y_error_avg, data_tuple_avg,
      deriv_sigmoid_avg)

```

^^^

^^^

^^^

the values fed to the backprop function for its three arguments are averaged over the training samples in the batch. This in keeping with the spirit of SGD that calls for averaging the information retained in the forward propagation over the samples in a batch.

See Slide 59 of my Week 3 slides for the math of back propagation for the One-Neuron network.

```

"""
input_vars = self.independent_vars
input_vars_to_param_map = self.var_to_var_param[
    self.output_vars[0]]
param_to_vars_map = {param: var for var, param in
    input_vars_to_param_map.items()}
vals_for_input_vars_dict = dict(zip(input_vars,
    list(vals_for_input_vars)))
vals_for_learnable_params = self.
    vals_for_learnable_params
for i, param in enumerate(self.
    vals_for_learnable_params):
    ## Calculate the next step in the parameter
    hyperplane
    #          step = self.learning_rate *
    y_error * vals_for_input_vars_dict[
    input_vars[i]] * deriv_sigmoid
grad = y_error * vals_for_input_vars_dict[
    param_to_vars_map[param]] * deriv_sigmoid
m = self.beta1 * self.prev_m[param] + (1-self.
    .beta1) * grad
v = self.beta2 * self.prev_v[param] + (1-self.
    .beta2) * (grad ** 2)
step = self.learning_rate * ((self.prev_m[
    param]/(1-self.beta1 ** (i+1)))/np.sqrt((

```

```

        self.prev_v[param]/(1-self.beta1 ** (i+1))
    )+self.epsilon))
    ## Update the learnable parameters
    self.prev_m[param] = m
    self.prev_v[param] = v
    self.vals_for_learnable_params[param] += step
    grad = y_error * deriv_sigmoid
    self.prev_biasv = self.beta1 * self.prev_biasm +
    (1-self.beta1) * grad
    self.prev_biasm = self.beta2 * self.prev_biasv +
    (1-self.beta2) * (grad ** 2)
    self.bias -= self.learning_rate * (self.
    prev_biasm/(self.prev_biasv+self.epsilon))

cgpp = ComputationalGraphPrimerPlus(
    one_neuron_model = True,
    expressions = ['xw=ab*xa+bc*xb+cd*xc+ac*xd
    '],
    output_vars = ['xw'],
    dataset_size = 5000,
    # learning_rate = 1e-3,
    learning_rate = 5 * 1e-5,
    training_iterations = 40000,
    batch_size = 8,
    display_loss_how_often = 100,
    debug = True,
)
cgpa = ComputationalGraphPrimerAdam(
    one_neuron_model = True,
    expressions = ['xw=ab*xa+bc*xb+cd*xc+ac*xd
    '],
    output_vars = ['xw'],
    dataset_size = 5000,
    # learning_rate = 1e-3,
    learning_rate = 5 * 1e-5,
    training_iterations = 40000,
    batch_size = 8,
    display_loss_how_often = 100,
    debug = True,
)

cgpp.parse_expressions()
cgpa.parse_expressions()
#cgp.display_network1()

```

```

# cgpp.display_network2()
training_data = cgpp.gen_training_data()
loss = cgpp.run_training_loop_one_neuron_model(
    training_data, 0.0)
loss_plus = cgpp.run_training_loop_one_neuron_model(
    training_data, 0.99)
loss_adam = cgpp.run_training_loop_one_neuron_model(
    training_data, 0.9, 0.99)

plt.figure()
plt.ylabel('loss')
plt.xlabel('iterations')
plt.title('SGD+ vs SGD vs Adam Loss')
plt.plot(loss_adam, label = 'Adam Training Loss', color='g')
plt.plot(loss_plus, label = 'SGD+ Training Loss', color='r')
plt.plot(loss, label = 'SGD Training Loss', color='b')
plt.legend()
#plt.show()
plt.savefig("one_neuron_loss_alt.jpg")

```

---

multi neuron:

---

```


#!/usr/bin/env python

## multi_neuron_classifier.py

"""
The main point of this script is to demonstrate saving
the information during the
forward propagation of data through a neural network and
using that information for
backpropagating the loss and for updating the values for
the learnable parameters. The
script uses the following 4-2-1 network layout, with 4
nodes in the input layer, 2 in
the hidden layer and 1 in the output layer as shown below
:

```

input



To explain what information is stored during the forward pass and how that information is used during the backprop step, see the comment blocks associated with the functions

```

forward_prop_multi_neuron_model()
and
backprop_and_update_params_multi_neuron_model()

```

Both of these functions are called by the training function:

```

run_training_loop_multi_neuron_model()

"""
import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
import sys
sys.path.append("E:\ECE60146DL\hw3_new\
ComputationalGraphPrimer-1.1.2\
ComputationalGraphPrimer")

import sys,os,os.path
import numpy as np
import re
import operator
import math

```



```

import random
import torch
from collections import deque
import copy
import matplotlib.pyplot as plt
import networkx as nx

seed = 0
random.seed(seed)
np.random.seed(seed)

from ComputationalGraphPrimer import *

class ComputationalGraphPrimerPlus(
    ComputationalGraphPrimer):
    def run_training_loop_multi_neuron_model(self,
        training_data, momentum):
        self.gamma = momentum
        self.prev_grad = {param: 0 for param in self.
            learnable_params}
        self.prev_bias = [0 for _ in range(self.
            num_layers-1)]
    class DataLoader:
        """
        To understand the logic of the dataloader, it
        would help if you first understand how
        the training dataset is created. Search for
        the following function in this file:

            gen_training_data(self)

        As you will see in the implementation code
        for this method, the training dataset
        consists of a Python dict with two keys, 0
        and 1, the former points to a list of
        all Class 0 samples and the latter to a list
        of all Class 1 samples. In each list,
        the data samples are drawn from a multi-
        dimensional Gaussian distribution. The
        two
        classes have different means and variances.
        The dimensionality of each data sample
        is set by the number of nodes in the input
        layer of the neural network.

        The data loader's job is to construct a batch

```

of samples drawn randomly from the two lists mentioned above. And it must also associate the class label with each sample separately.

```
"""
def __init__(self, training_data, batch_size):
    :
    self.training_data = training_data
    self.batch_size = batch_size
    self.class_0_samples = [(item, 0) for
        item in
            self.
                training_data
                [0]] ##
                Associate
                label 0 with
                each sample
    self.class_1_samples = [(item, 1) for
        item in
            self.
                training_data
                [1]] ##
                Associate
                label 1 with
                each sample

def __len__(self):
    return len(self.training_data[0]) + len(
        self.training_data[1])

def _getitem(self):
    cointoss = random.choice([0, 1]) ## When
        a batch is created by getbatch(), we
        want the
    ## samples to be chosen randomly from
        the two lists
    if cointoss == 0:
        return random.choice(self.
            class_0_samples)
    else:
        return random.choice(self.
            class_1_samples)

def getbatch(self):
    batch_data, batch_labels = [], [] ##
```

```

        First list for samples, the second for
        labels
        maxval = 0.0  ## For approximate batch
        data normalization
        for _ in range(self.batch_size):
            item = self._getitem()
            if np.max(item[0]) > maxval:
                maxval = np.max(item[0])
            batch_data.append(item[0])
            batch_labels.append(item[1])
        batch_data = [item / maxval for item in
            batch_data]  ## Normalize batch data
        batch = [batch_data, batch_labels]
        return batch

"""
The training loop must first initialize the
    learnable parameters. Remember, these are the
symbolic names in your input expressions for the
    neural layer that do not begin with the
letter 'x'. In this case, we are initializing
    with random numbers from a uniform
    distribution
over the interval (0,1).
"""
self.vals_for_learnable_params = {param: random.
    uniform(0, 1) for param in self.
    learnable_params}

self.bias = [random.uniform(0, 1) for _ in
    range(self.num_layers - 1)]  ##
    Adding the bias to each layer
    improves
## class discrimination. We initialize it
## to a random number.

data_loader = DataLoader(training_data,
    batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_iterations = 0.0  ## Average the
    loss over iterations for printing out
## every N iterations during the training loop
.
for i in range(self.training_iterations):
    data = data_loader.getbatch()

```

```

data_tuples = data[0]
class_labels = data[1]
self.forward_prop_multi_neuron_model(
    data_tuples) ## FORW PROP works by side-
                effect
predicted_labels_for_batch = self.
    forw_prop_vals_at_layers[
        self.num_layers - 1] ## Predictions from
                            FORW PROP
y_preds = [item for sublist in
    predicted_labels_for_batch for item in
        sublist] ## Get numeric vals for
                predictions
loss = sum([(abs(class_labels[i] - y_preds[i]
    ])) ** 2 for i in
        range(len(class_labels))]) ##
                            Calculate loss for batch
loss_avg = loss / float(len(class_labels))
## Average the loss over batch
avg_loss_over_iterations += loss_avg ## Add
to Average loss over iterations
if i % (self.display_loss_how_often) == 0:
    avg_loss_over_iterations /= self.
        display_loss_how_often
    loss_running_record.append(
        avg_loss_over_iterations)
    print("[iter=%d] loss = %.4f" % (i + 1,
        avg_loss_over_iterations)) ## Display
        avg loss
    avg_loss_over_iterations = 0.0 ## Re-
        initialize avg-over-iterations loss
y_errors = list(map(operator.sub,
    class_labels, y_preds))
y_error_avg = sum(y_errors) / float(len(
    class_labels))
self.
    backprop_and_update_params_multi_neuron_model(
        (y_error_avg, class_labels)) ## BACKPROP
        loss
return loss_running_record

```

```

def backprop_and_update_params_multi_neuron_model(
    self, y_error, class_labels):
    """

```

First note that loop index variable '`back_layer_index`' starts with the index of the last layer. For the 3-layer example shown for 'forward', `back_layer_index` starts with a value of 2, its next value is 1, and that's it.

Stochastic Gradient Gradient calls for the backpropagated loss to be averaged over the samples in a batch. To explain how this averaging is carried out by the `backprop` function, consider the last node on the example shown in the `forward()` function above. Standing at the node, we look at the 'input' values stored in the variable "`input_vals`". Assuming a batch size of 8, this will be list of lists. Each of the inner lists will have two values for the two nodes in the hidden layer. And there will be 8 of these for the 8 elements of the batch. We average these values 'input\_vals' and store those in the variable "`input_vals_avg`". Next we must carry out the same batch-based averaging for the partial derivatives stored in the variable "`deriv_sigmoid`".

Pay attention to the variable '`vars_in_layer`'. These store the node variables in the current layer during backpropagation. Since `back_layer_index` starts with a value of 2, the variable '`vars_in_layer`' will have just the single node for the example shown for `forward()`. With respect to what is stored in `vars_in_layer`, the variables stored in '`input_vars_to_layer`' correspond to the input layer with respect to the current layer.

```
"""
# backproped prediction error:
pred_err_backproped_at_layers = {i: [] for i in
    range(1, self.num_layers - 1)}
pred_err_backproped_at_layers[self.num_layers -
    1] = [y_error]
for back_layer_index in reversed(range(1, self.
    num_layers)):
```

```

input_vals = self.forw_prop_vals_at_layers[
    back_layer_index - 1]
input_vals_avg = [sum(x) for x in zip(*
    input_vals)]
input_vals_avg = list(map(operator.truediv,
    input_vals_avg, [float(len(class_labels))
    * len(class_labels)]))
deriv_sigmoid = self.gradient_vals_for_layers
    [back_layer_index]
deriv_sigmoid_avg = [sum(x) for x in zip(*
    deriv_sigmoid)]
deriv_sigmoid_avg = list(map(operator.truediv
    , deriv_sigmoid_avg,
                                [float(len(
                                    class_labels)
                                )] * len(
                                    class_labels)
                                ))
vars_in_layer = self.layer_vars[
    back_layer_index] ## a list like ['xo']
vars_in_next_layer_back = self.layer_vars[
    back_layer_index - 1] ## a list like ['xw
    ', 'xz']

layer_params = self.layer_params[
    back_layer_index]
## note that layer_params are stored in a
dict like
## {1: [['ap', 'aq', 'ar', 'as'], ['bp',
    'bq', 'br', 'bs']], 2: [['cp', 'cq']]}
## "layer_params[idx]" is a list of lists for
the link weights in layer whose output
nodes are in layer "idx"
transposed_layer_params = list(zip(*
    layer_params)) ## creating a transpose of
the link matrix

backproped_error = [None] * len(
    vars_in_next_layer_back)
for k, varr in enumerate(
    vars_in_next_layer_back):
    for j, var2 in enumerate(vars_in_layer):
        backproped_error[k] = sum([self.
            vals_for_learnable_params[
                transposed_layer_params[k][i]] *
                pred_err_backproped_at_layers

```

```

[
    back_layer_index
][i]
for i in
    range(
        len(
            vars_in_layer
        ))))

#

    deriv_sigmoid_avg[i] for i in range(len(
        vars_in_layer)))
pred_err_backproped_at_layers[
    back_layer_index - 1] = backproped_error
input_vars_to_layer = self.layer_vars[
    back_layer_index - 1]
for j, var in enumerate(vars_in_layer):
    layer_params = self.layer_params[
        back_layer_index][j]
    ## Regarding the parameter update loop
    that follows, see the Slides 74
    through 77 of my Week 3
    ## lecture slides for how the parameters
    are updated using the partial
    derivatives stored away
    ## during forward propagation of data.
    The theory underlying these
    calculations is presented
    ## in Slides 68 through 71.
    for i, param in enumerate(layer_params):
        gradient_of_loss_for_param =
            input_vals_avg[i] *
            pred_err_backproped_at_layers[
                back_layer_index][j]
        grad = gradient_of_loss_for_param *
            deriv_sigmoid_avg[j]
        self.prev_grad[param] = self.
            learning_rate * grad + self.gamma
            * self.prev_grad[param]
        self.vals_for_learnable_params[param]
            += self.prev_grad[param]
self.prev_bias[back_layer_index - 1] = self.
    learning_rate * sum(
        pred_err_backproped_at_layers[
            back_layer_index]) * sum(deriv_sigmoid_avg
    )/len(deriv_sigmoid_avg) + self.gamma *

```

```

        self.prev_bias[back_layer_index-1]
self.bias[back_layer_index - 1] += self.
prev_bias[back_layer_index - 1]

```

```

class ComputationalGraphPrimerAdam(
    ComputationalGraphPrimer):
    def run_training_loop_multi_neuron_model(self,
        training_data, beta1, beta2):
        self.beta1 = beta1
        self.beta2 = beta2
        self.epsilon = 1e-8
        self.prev_m = {param: 0 for param in self.
            learnable_params}
        self.prev_v = {param: 0 for param in self.
            learnable_params}
        self.prev_biasv = [0 for _ in range(self.
            num_layers-1)]
        self.prev_biasm = [0 for _ in range(self.
            num_layers - 1)]
class DataLoader:
    """

```

To understand the logic of the dataloader, it would help if you first understand how the training dataset is created. Search for the following function in this file:

```
gen_training_data(self)
```

As you will see in the implementation code for this method, the training dataset consists of a Python dict with two keys, 0 and 1, the former points to a list of all Class 0 samples and the latter to a list of all Class 1 samples. In each list, the data samples are drawn from a multi-dimensional Gaussian distribution. The two classes have different means and variances. The dimensionality of each data sample is set by the number of nodes in the input layer of the neural network.

The data loader's job is to construct a batch of samples drawn randomly from the two



lists mentioned above. And it must also  
 associate the class label with each sample  
 separately.  
 """

```
def __init__(self, training_data, batch_size):
    :
    self.training_data = training_data
    self.batch_size = batch_size
    self.class_0_samples = [(item, 0) for
        item in
            self.
                training_data
                [0]] ##
                Associate
                label 0 with
                each sample
    self.class_1_samples = [(item, 1) for
        item in
            self.
                training_data
                [1]] ##
                Associate
                label 1 with
                each sample

def __len__(self):
    return len(self.training_data[0]) + len(
        self.training_data[1])

def _getitem(self):
    cointoss = random.choice([0, 1]) ## When
        a batch is created by getbatch(), we
        want the
    ## samples to be chosen randomly from
        the two lists
    if cointoss == 0:
        return random.choice(self.
            class_0_samples)
    else:
        return random.choice(self.
            class_1_samples)

def getbatch(self):
    batch_data, batch_labels = [], [] ##
        First list for samples, the second for
```

```

        labels
        maxval = 0.0  ## For approximate batch
        data normalization
        for _ in range(self.batch_size):
            item = self._getitem()
            if np.max(item[0]) > maxval:
                maxval = np.max(item[0])
            batch_data.append(item[0])
            batch_labels.append(item[1])
        batch_data = [item / maxval for item in
            batch_data]  ## Normalize batch data
        batch = [batch_data, batch_labels]
        return batch

"""
The training loop must first initialize the
    learnable parameters. Remember, these are the
symbolic names in your input expressions for the
    neural layer that do not begin with the
letter 'x'. In this case, we are initializing
    with random numbers from a uniform
    distribution
over the interval (0,1).
"""

self.vals_for_learnable_params = {param: random.
    uniform(0, 1) for param in self.
    learnable_params}

self.bias = [random.uniform(0, 1) for _ in
    range(self.num_layers - 1)]  ##
    Adding the bias to each layer
    improves
## class discrimination. We initialize it
## to a random number.

data_loader = DataLoader(training_data,
    batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_iterations = 0.0  ## Average the
    loss over iterations for printing out
## every N iterations during the training loop
.
for i in range(self.training_iterations):
    data = data_loader.getbatch()
    data_tuples = data[0]

```

```

class_labels = data[1]
self.forward_prop_multi_neuron_model(
    data_tuples) ## FORW PROP works by side-
                effect
predicted_labels_for_batch = self.
    forw_prop_vals_at_layers[
        self.num_layers - 1] ## Predictions from
                            FORW PROP
y_preds = [item for sublist in
    predicted_labels_for_batch for item in
        sublist] ## Get numeric vals for
                predictions
loss = sum([(abs(class_labels[i] - y_preds[i
    ])) ** 2 for i in
        range(len(class_labels))]) ##
                            Calculate loss for batch
loss_avg = loss / float(len(class_labels))
## Average the loss over batch
avg_loss_over_iterations += loss_avg ## Add
to Average loss over iterations
if i % (self.display_loss_how_often) == 0:
    avg_loss_over_iterations /= self.
        display_loss_how_often
    loss_running_record.append(
        avg_loss_over_iterations)
    print("[iter=%d] loss = %.4f" % (i + 1,
        avg_loss_over_iterations)) ## Display
        avg loss
    avg_loss_over_iterations = 0.0 ## Re-
        initialize avg-over-iterations loss
y_errors = list(map(operator.sub,
    class_labels, y_preds))
y_error_avg = sum(y_errors) / float(len(
    class_labels))
self.
    backprop_and_update_params_multi_neuron_model
        (y_error_avg, class_labels) ## BACKPROP
        loss
return loss_running_record

```

```

def backprop_and_update_params_multi_neuron_model(
    self, y_error, class_labels):
    """

```

First note that loop index variable '

`back_layer_index` starts with the index of the last layer. For the 3-layer example shown for 'forward', `back_layer_index` starts with a value of 2, its next value is 1, and that's it.

Stochastic Gradient Gradient calls for the backpropagated loss to be averaged over the samples in a batch. To explain how this averaging is carried out by the `backprop` function, consider the last node on the example shown in the `forward()` function above. Standing at the node, we look at the 'input' values stored in the variable `"input_vals"`. Assuming a batch size of 8, this will be list of lists. Each of the inner lists will have two values for the two nodes in the hidden layer. And there will be 8 of these for the 8 elements of the batch. We average these values 'input\_vals' and store those in the variable `"input_vals_avg"`. Next we must carry out the same batch-based averaging for the partial derivatives stored in the variable `"deriv_sigmoid"`.

Pay attention to the variable `'vars_in_layer'`. These store the node variables in the current layer during backpropagation. Since `back_layer_index` starts with a value of 2, the variable `'vars_in_layer'` will have just the single node for the example shown for `forward()`. With respect to what is stored in `vars_in_layer`, the variables stored in `'input_vars_to_layer'` correspond to the input layer with respect to the current layer.

```
"""
# backproped prediction error:
pred_err_backproped_at_layers = {i: [] for i in
    range(1, self.num_layers - 1)}
pred_err_backproped_at_layers[self.num_layers -
    1] = [y_error]
for back_layer_index in reversed(range(1, self.
    num_layers)):
    input_vals = self.forw_prop_vals_at_layers[
```

```

        back_layer_index - 1]
input_vals_avg = [sum(x) for x in zip(*
    input_vals)]
input_vals_avg = list(map(operator.truediv,
    input_vals_avg, [float(len(class_labels))
    * len(class_labels)]))
deriv_sigmoid = self.gradient_vals_for_layers
    [back_layer_index]
deriv_sigmoid_avg = [sum(x) for x in zip(*
    deriv_sigmoid)]
deriv_sigmoid_avg = list(map(operator.truediv
    , deriv_sigmoid_avg,
                                [float(len(
                                    class_labels)
                                )] * len(
                                    class_labels)
                                ))
vars_in_layer = self.layer_vars[
    back_layer_index] ## a list like ['xo']
vars_in_next_layer_back = self.layer_vars[
    back_layer_index - 1] ## a list like ['xw
    ', 'xz']

layer_params = self.layer_params[
    back_layer_index]
## note that layer_params are stored in a
dict like
## {1: [['ap', 'aq', 'ar', 'as'], ['bp',
    'bq', 'br', 'bs']], 2: [['cp', 'cq']]}
## "layer_params[idx]" is a list of lists for
the link weights in layer whose output
nodes are in layer "idx"
transposed_layer_params = list(zip(*
    layer_params)) ## creating a transpose of
the link matrix

backproped_error = [None] * len(
    vars_in_next_layer_back)
for k, varr in enumerate(
    vars_in_next_layer_back):
    for j, var2 in enumerate(vars_in_layer):
        backproped_error[k] = sum([self.
            vals_for_learnable_params[
                transposed_layer_params[k][i]] *
                pred_err_backproped_at_layers
            [

```

```

back_layer_index
][i]
for i in
    range(
        len(
            vars_in_layer
        )))
#
    deriv_sigmoid_avg[i] for i in range(len(
        vars_in_layer)))
pred_err_backproped_at_layers[
    back_layer_index - 1] = backproped_error
input_vars_to_layer = self.layer_vars[
    back_layer_index - 1]
for j, var in enumerate(vars_in_layer):
    layer_params = self.layer_params[
        back_layer_index][j]
    ## Regarding the parameter update loop
    ## that follows, see the Slides 74
    ## through 77 of my Week 3
    ## lecture slides for how the parameters
    ## are updated using the partial
    ## derivatives stored away
    ## during forward propagation of data.
    ## The theory underlying these
    ## calculations is presented
    ## in Slides 68 through 71.
    for i, param in enumerate(layer_params):
        gradient_of_loss_for_param =
            input_vals_avg[i] *
            pred_err_backproped_at_layers[
                back_layer_index][j]
        grad = gradient_of_loss_for_param *
            deriv_sigmoid_avg[j]
        m = self.beta1 * self.prev_m[param] +
            (1 - self.beta1) * grad
        v = self.beta2 * self.prev_v[param] +
            (1 - self.beta2) * (grad ** 2)
        step = self.learning_rate * ((self.
            prev_m[param] / (1 - self.beta1 **
                (i + 1))) / np.sqrt((self.prev_v[
                    param] / (1 - self.beta1 ** (i +
                        1))) + self.epsilon))
        self.prev_m[param] = m
        self.prev_v[param] = v

```



```

        'xo=cp*xw+cq*xz'],
        output_vars = ['xo'],
        dataset_size = 5000,
        # learning_rate = 1e-3,
        learning_rate = 5 * 1e-5,
        training_iterations = 40000,
        batch_size = 8,
        display_loss_how_often = 100,
        debug = True,
    )

    cgp.parse_multi_layer_expressions()
    cgpa.parse_multi_layer_expressions()

    #cgp.display_network1()
    # cgpa.display_network2()

    training_data = cgp.gen_training_data()

    loss = cgp.run_training_loop_multi_neuron_model(
        training_data, 0.0 )
    loss_plus = cgp.run_training_loop_multi_neuron_model(
        training_data, 0.99)
    loss_adam = cgpa.run_training_loop_multi_neuron_model(
        training_data, 0.9, 0.99)
    plt.figure()
    plt.ylabel('loss')
    plt.xlabel('iterations')
    plt.title('SGD+ vs SGD vs adam Loss')
    plt.plot(loss_adam, label = 'Adam Training Loss', color='
        g')
    plt.plot(loss_plus, label = 'SGD+ Training Loss', color='
        r')
    plt.plot(loss, label = 'SGD Training Loss', color='b')
    plt.legend()
    #plt.show()
    plt.savefig("multi_neuron_loss_alt2.jpg")

```

---